

Zhiyue Qiu

Algorithm:

I use 'push' algorithm and parallel it.

This algorithm and implementation detail for each process goes like this:

Main Algorithm:

For each iteration:

 initialize send and receive buffer

 for each vertex:

 calculate the 'push' value

 for each adjacent vertex:

 store the 'push' key in send buffer

 store the 'push' value in send buffer

 MPI_Scatterv: send data to each processor

 for each received chunk of data:

 update key value

 for each PR:

 calculate norm_changed

 MPI_Reduce on norm_changed and determine if terminate the loop

End

Detailed implementation:

Myid == 0:

 read in data, Send chunks to each processor. Chunk size is determined by NNZ of each vertex. Goal is to make sure each processor have similar number of NNZ.

 NNZ might close to E/p , which is the data, each processor should get. Also calculate and broadcast block_size, total number of vertices, edges to each processor.

 Each processor allocate sending and receiving buffer. The size should be the num it is to receive and send. Each processor should know, how much I need to send to and how much I need to receive from. Also calculate the displacement for MPI_Scatterv, which can send to each processor different size of data.

Do the algorithm above

Zhiyue Qiu

Complexity analysis:

For each iteration:

- initialize send and receive buffer $\rightarrow O(E/p)$
- for each vertex: $\rightarrow O(V/p + E/p)$
 - calculate the 'push' value
 - for each adjacent vertex:
 - store the 'push' key in send buffer
 - store the 'push' value in send buffer

MPI_Scatterv: send data to each processor

$\rightarrow p * \text{one to all communication: } p * O(\log p * T_s + (p-1)/p * E * T_w)$

- for each received chunk of data: $\rightarrow O(E/p)$
 - update key value

- for each PR: $\rightarrow O(V/p)$
 - calculate norm_changed

MPI_Reduce on norm_changed and determine if terminate the loop
 $\rightarrow O(\log p * T_s + T_w)$

End

Serial complexity: (in each iteration) $O(V+E)$

$p * T_p = O(E) + O(V) + p * \log p * T_s + (p-1) * E * T_w$

Overhead: Mainly from scatter

$p * \log p * T_s + (p-1) * E * T_w$

It behaves better on A.graph, because A.graph is well formed, every chunk of nodes only send data to another chunk, so the overhead might become something like $p * T_s + E * T_w$.

B.graph and live.graph, worse on B, Better on live. The reason is that based on my observation, live.graph is not that mess. But B is purely random.

So A is the best case, every processor only need to communicate with one other processor. B is the worst, the real scatter p times. The live.graph between them.

Timing result

A.graph

```
mpirun -np 1 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 1.904s
```

```
mpirun -np 2 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 1.094s
```

```
mpirun -np 4 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 0.760s
```

```
mpirun -np 8 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 0.580s
```

```
mpirun -np 16 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 0.484s
```

```
mpirun -np 32 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
A.graph A-parallel.out
```

```
Number of iterations: 38 average time: 0.477s
```

B.graph

```
mpirun -np 1 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

```
Number of iterations: 8 average time: 4.304s
```

```
mpirun -np 2 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

```
Number of iterations: 8 average time: 1.522s
```

```
mpirun -np 4 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

```
Number of iterations: 8 average time: 0.972s
```

```
mpirun -np 8 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

```
Number of iterations: 8 average time: 0.694s
```

```
mpirun -np 16 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

```
Number of iterations: 8 average time: 0.577s
```

```
mpirun -np 32 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/
```

```
B.graph B-parallel.out
```

Zhiyue Qiu

Number of iterations: 8 average time: 0.628s

live.graph

```
mpirun -np 1 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 2.208s

```
mpirun -np 2 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 1.338s

```
mpirun -np 4 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 0.996s

```
mpirun -np 8 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 0.806s

```
mpirun -np 16 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 0.664s

```
mpirun -np 32 ./pagerank /export/scratch/CSCI5451_S18/assignment-4/  
live-journal.graph live-parallel.out
```

Number of iterations: 62 average time: 0.799s